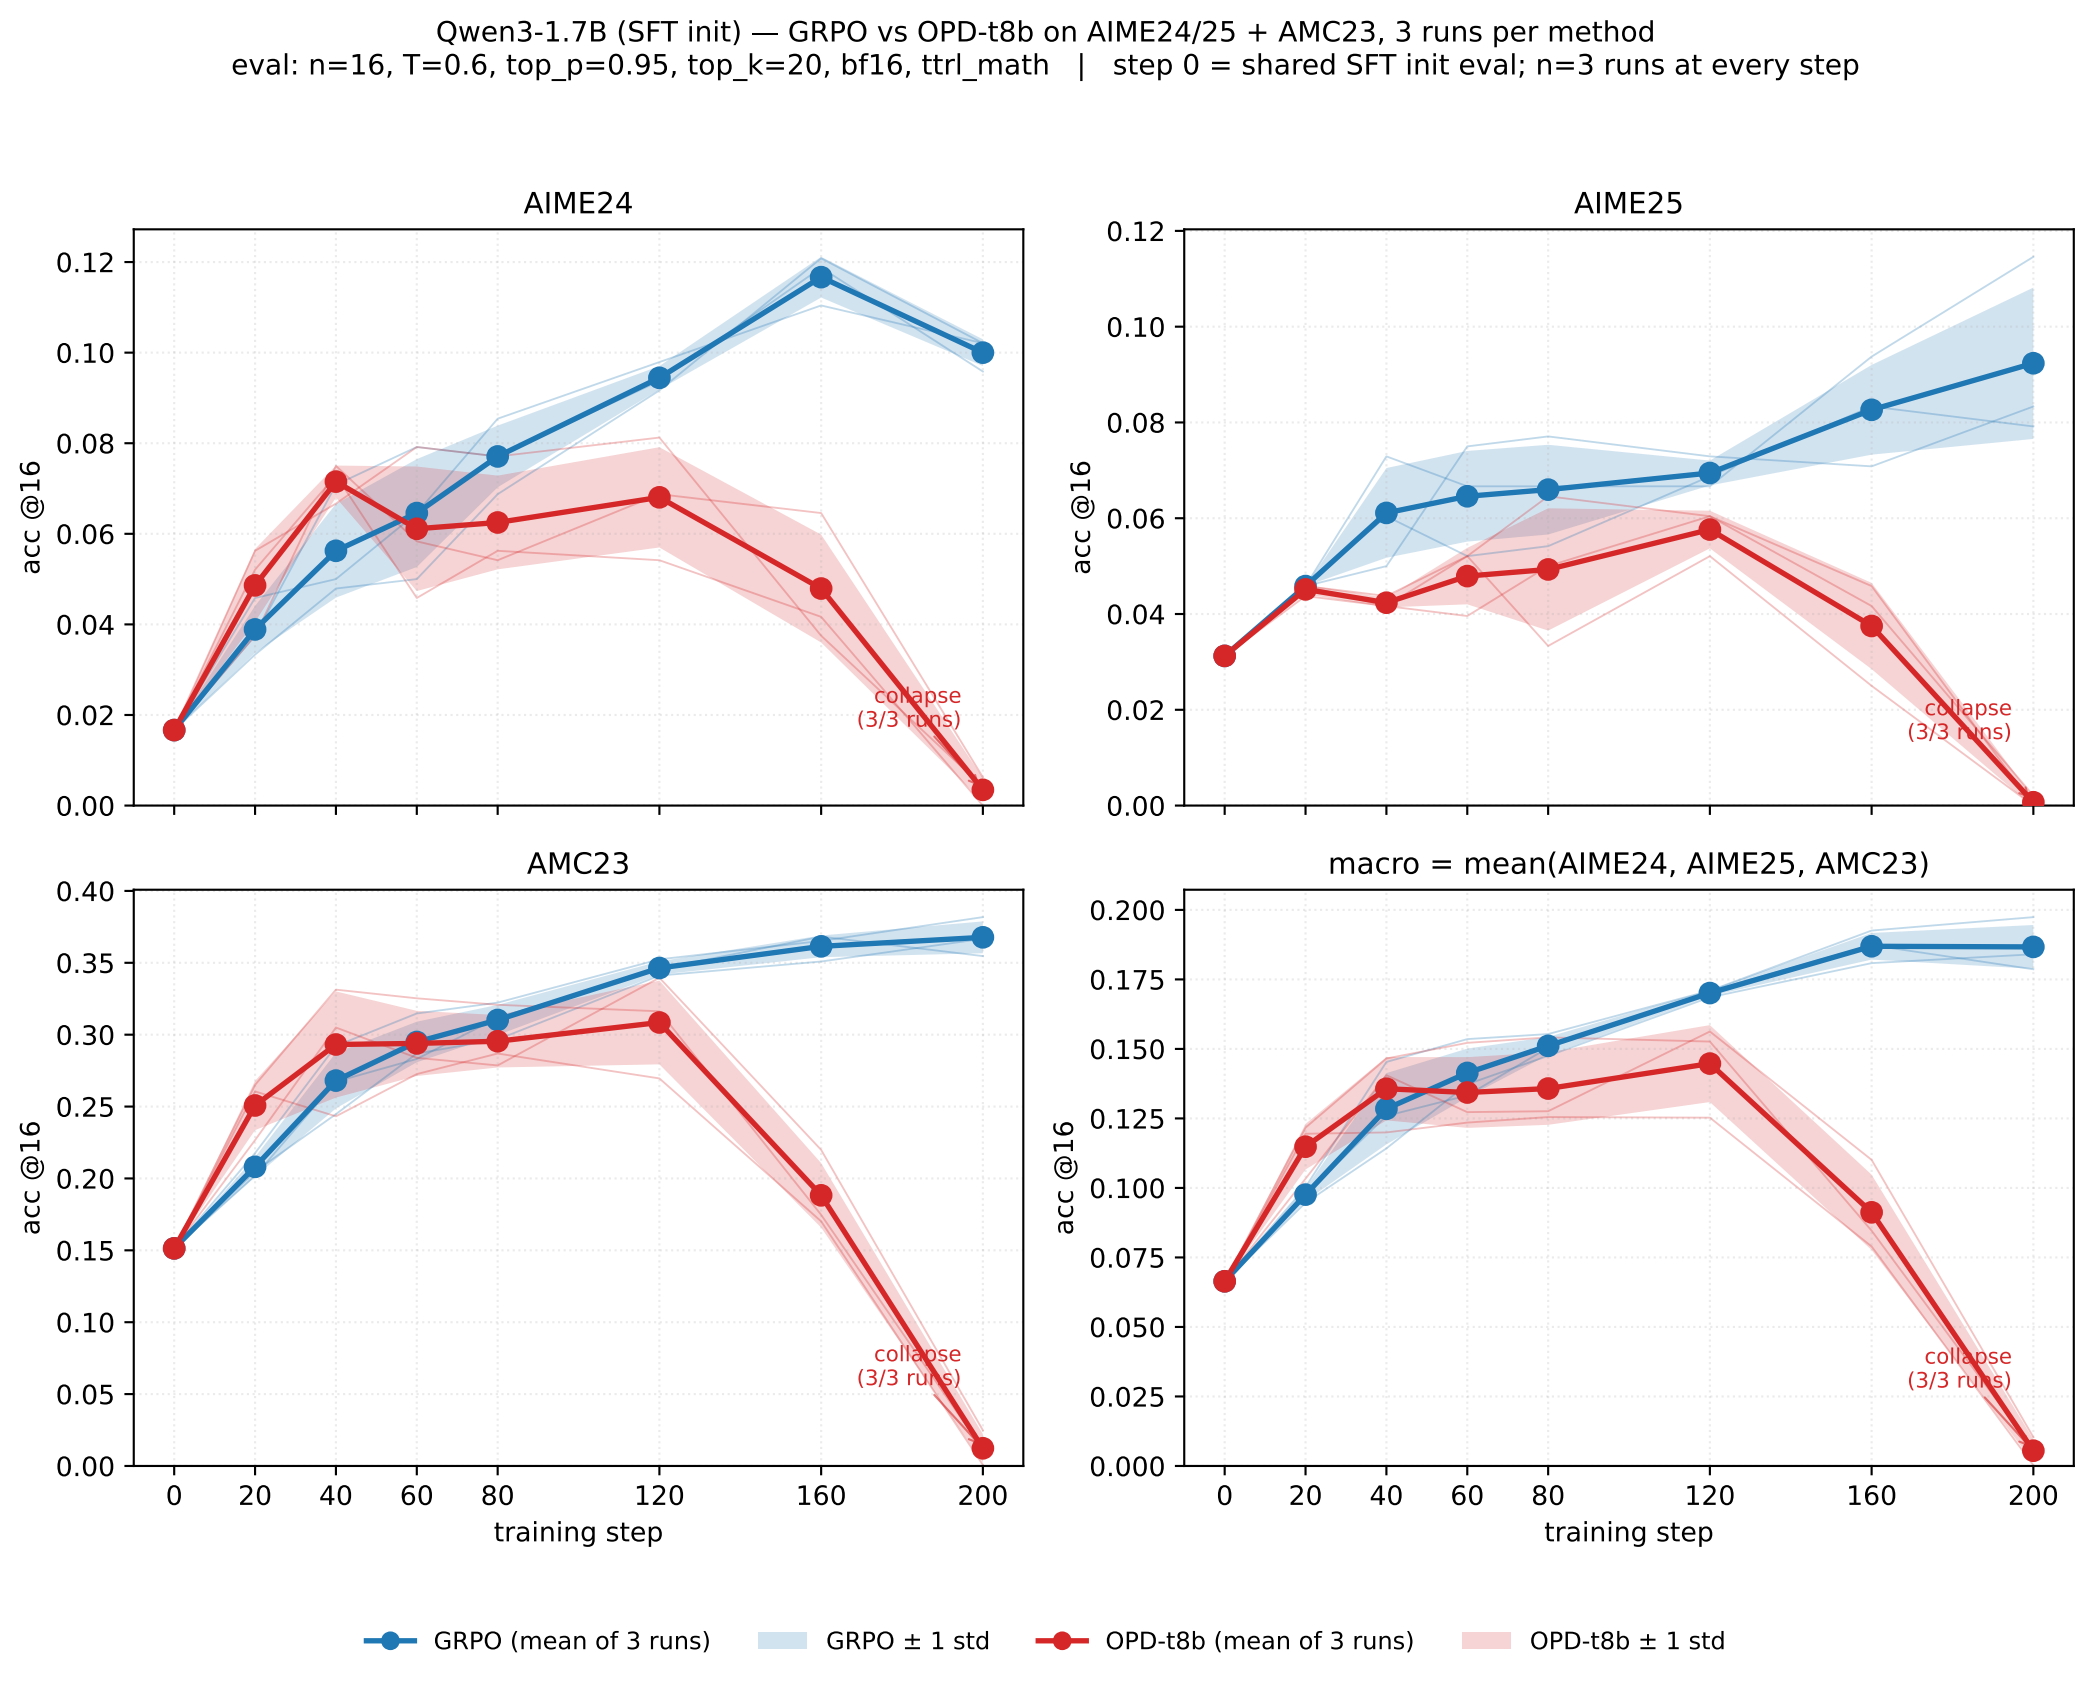
Qwen3-1.7B (SFT init) — GRPO vs OPD-t8b on AIME24/25 + AMC23, 3 runs per method
eval: n=16, T=0.6, top_p=0.95, top_k=20, bf16, ttrl_math | step 0 = shared SFT init eval; n=3 runs at every step
AIME24
AIME25
AMC23
macro = mean(AIME24, AIME25, AMC23)
acc @16
training step
collapse (3/3 runs)
GRPO (mean of 3 runs)
GRPO ± 1 std
OPD-t8b (mean of 3 runs)
OPD-t8b ± 1 std